

Step-by-Step Explanation

This code implements a document scanner that takes an image of a document (e.g., a piece of paper) and processes it to produce a clean, flat, scanned version. Below is a detailed explanation of each step, designed to be clear for a student new to image processing.

How It All Comes Together

The code takes an image of a document (e.g., a piece of paper on a table), identifies its boundaries, corrects its perspective to make it look flat, and enhances it to resemble a scanned document. Each step builds on the previous one:

- **Load and Resize:** Prepare the image for processing.
- **Grayscale and Edge Detection:** Simplify the image and find its edges.
- **Contour Detection:** Identify the document's shape.
- **Perspective Transform:** Flatten the document.
- **Thresholding:** Enhance the document for clarity.

Step 1: Setting Up the Environment

- **What's Happening:** Before running the code, you need to install the required Python libraries: `opencv-python` (for image processing), `imutils` (for resizing images), `scikit-image` (for thresholding), and `numpy` (for array operations).
- **Code:**

```
pip install opencv-python imutils scikit-image numpy
```

- **Explanation:** These libraries provide tools to manipulate images. OpenCV (`cv2`) is the main library for image processing, `imutils` simplifies tasks like resizing, `scikit-image` helps with advanced image filtering, and `numpy` handles numerical data like image pixel arrays.
- **Why It Matters:** Installing these libraries ensures you have the tools needed to process images effectively.

Step 2: Importing Libraries

- **What's Happening:** The code imports the necessary libraries and defines a custom `order_points` function to arrange the document's corner points.
- **Python Code:**

```
import cv2

import numpy as np

import imutils

from skimage.filters import threshold_local

from scipy.spatial import distance as dist
```

- **Explanation:**

- cv2 is OpenCV for image processing tasks like reading images, converting colors, and detecting edges.
- numpy (np) handles arrays, which represent images in this context.
- imutils provides a simple function to resize images.
- threshold_local from scikit-image applies adaptive thresholding to enhance the document's text.
- scipy.spatial.distance is used in the order_points function to calculate distances between points.

- **Why It Matters:** These imports give you access to all the functions needed to process the image step-by-step.

Step 3: Loading and Resizing the Image

- **What's Happening:** The code loads an image (e.g., sample.jpg) and resizes it to a height of 500 pixels while maintaining the aspect ratio.
- **Python Code:**

```
image_path = 'images/sample.jpg'

original_img = cv2.imread(image_path)

copy = original_img.copy()

ratio = original_img.shape[0] / 500.0

img_resize = imutils.resize(original_img, height=500)
```

- **Explanation:**

- cv2.imread loads the image from the specified path into a NumPy array.
- copy creates a duplicate of the original image for later use in perspective transformation.
- original_img.shape[0] gets the image's height in pixels. Dividing by 500 calculates the ratio to maintain the aspect ratio.
- imutils.resize resizes the image to a height of 500 pixels, making it easier to process while keeping proportions intact.

- **Why It Matters:** Resizing reduces computational load and standardizes the image size for consistent processing. The copy preserves the original image for later steps.

Step 4: Converting to Grayscale

- **What's Happening:** The resized image is converted from color (RGB) to grayscale.
- **Python Code:**

```
gray_image = cv2.cvtColor(img_resize, cv2.COLOR_BGR2GRAY)
```

```
cv2.imshow("Grayscale Image", gray_image)
```

```
cv2.waitKey(0)
```

- **Explanation:**
 - `cv2.cvtColor` converts the image from BGR (OpenCV's default color format) to grayscale (`cv2.COLOR_BGR2GRAY`), reducing it to a single channel (intensity values).
 - Grayscale images are simpler to process because they have one value per pixel (0–255) instead of three (red, green, blue).
 - `cv2.imshow` displays the grayscale image, and `cv2.waitKey(0)` pauses execution until a key is pressed.
- **Why It Matters:** Grayscale simplifies subsequent steps like edge detection, as it reduces the complexity of the image data.

Step 5: Applying Gaussian Blur and Edge Detection

- **What's Happening:** The grayscale image is blurred to reduce noise, and then edges are detected using the Canny edge detector.
- **Python Code:**

```
blurred_image = cv2.GaussianBlur(gray_image, (5, 5), 0)
```

```
edged_image = cv2.Canny(blurred_image, 75, 200)
```

```
cv2.imshow("Edged Image", edged_image)
```

```
cv2.waitKey(0)
```

- **Explanation:**
 - `cv2.GaussianBlur` applies a 5x5 Gaussian blur to smooth the image, reducing noise that could interfere with edge detection.
 - `cv2.Canny` detects edges by identifying areas with significant intensity changes. The parameters 75 and 200 are thresholds for edge strength.
 - The result is a binary image (black background, white edges) highlighting the document's boundaries.

- The edged image is displayed for visualization.
- **Why It Matters:** Edge detection helps identify the document's outline, which is crucial for finding its shape in the next step.

Step 6: Finding Contours and Identifying the Document

- **What's Happening:** The code finds contours (outlines) in the edged image, selects the largest ones, and identifies the document by looking for a quadrilateral (four-sided shape).
- **Python Code:**

```
cnts = cv2.findContours(edged_image, cv2.RETR_LIST, cv2.CHAIN_APPROX_SIMPLE)
```

```
cnts = cnts[0] if len(cnts) == 2 else cnts[1]
```

```
cnts = sorted(cnts, key=cv2.contourArea, reverse=True)[:5]
```

```
doc = None
```

```
for c in cnts:
```

```
    peri = cv2.arcLength(c, True)
```

```
    approx = cv2.approxPolyDP(c, 0.02 * peri, True)
```

```
    if len(approx) == 4:
```

```
        doc = approx
```

```
    break
```

- **Explanation:**
 - `cv2.findContours` detects contours in the edged image. `cv2.RETR_LIST` retrieves all contours, and `cv2.CHAIN_APPROX_SIMPLE` simplifies their representation.
 - The code handles OpenCV version differences by checking the length of the returned tuple.
 - Contours are sorted by area (`cv2.contourArea`), and the top 5 largest are kept to focus on significant shapes.
 - For each contour, `cv2.arcLength` calculates its perimeter, and `cv2.approxPolyDP` approximates it to a simpler polygon. The `0.02 * peri` parameter controls the approximation accuracy.
 - If a contour has 4 sides (`len(approx) == 4`), it's likely the document, so it's stored in `doc` and the loop breaks.
- **Why It Matters:** The document is typically a rectangle, so finding a four-sided contour helps isolate it from other shapes in the image.

Step 7: Visualizing the Document's Corners

- **What's Happening:** The code draws red circles at the detected document's corners to visualize them.
- **Python Code:**

if doc is not None:

```
points = []  
  
for point in doc:  
    tuple_point = tuple(point[0])  
    cv2.circle(img_resize, tuple_point, 5, (0, 0, 255), -1)  
    points.append(tuple_point)  
  
cv2.imshow("Corner Points", img_resize)  
cv2.waitKey(0)
```

- **Explanation:**
 - If a document contour (doc) is found, the code loops through its 4 points.
 - cv2.circle draws a red circle ((0, 0, 255) in BGR, with radius 5) at each corner on the resized RGB image.
 - The points are stored in a list for potential use and displayed for verification.
- **Why It Matters:** Visualizing the corners confirms that the correct document shape has been detected before proceeding.

Step 8: Ordering the Corner Points

- **What's Happening:** The order_points function arranges the document's 4 corner points in a consistent order: top-left, top-right, bottom-right, bottom-left.
- **Python Code:**

```
def order_points(pts):  
    rect = np.zeros((4, 2), dtype="float32")  
  
    s = pts.sum(axis=1)  
    rect[0] = pts[np.argmin(s)]  
    rect[2] = pts[np.argmax(s)]  
  
    diff = np.diff(pts, axis=1)  
    rect[1] = pts[np.argmin(diff)]
```

```
rect[3] = pts[np.argmax(diff)]
```

```
return rect
```

- **Explanation:**
 - pts is an array of 4 points (x, y coordinates).
 - The function sums $x + y$ for each point. The point with the smallest sum is top-left, and the largest sum is bottom-right.
 - It computes the difference $(y - x)$ for each point. The smallest difference indicates top-right, and the largest difference indicates bottom-left.
 - The ordered points are returned as a 4x2 array.
- **Why It Matters:** Consistent ordering is necessary for the perspective transform to map the document correctly.

Step 9: Applying Perspective Transform

- **What's Happening:** The code uses the `perspective_transform` function to flatten the document into a top-down view.
- **Python Code:**

```
def perspective_transform(image, pts, ratio):  
    rect = order_points(pts)  
    (tl, tr, br, bl) = rect  
    widthA = np.sqrt(((br[0] - bl[0]) ** 2) + ((br[1] - bl[1]) ** 2))  
    widthB = np.sqrt(((tr[0] - tl[0]) ** 2) + ((tr[1] - tl[1]) ** 2))  
    maxWidth = max(int(widthA), int(widthB))  
    heightA = np.sqrt(((tr[0] - br[0]) ** 2) + ((tr[1] - br[1]) ** 2))  
    heightB = np.sqrt(((tl[0] - bl[0]) ** 2) + ((tl[1] - bl[1]) ** 2))  
    maxHeight = max(int(heightA), int(heightB))  
    dst = np.array([[0, 0], [maxWidth - 1, 0], [maxWidth - 1, maxHeight - 1], [0, maxHeight - 1]],  
dtype="float32")  
    M = cv2.getPerspectiveTransform(rect, dst)  
    warped = cv2.warpPerspective(image, M, (maxWidth, maxHeight))  
    return warped
```

- **Explanation:**
 - The function calls `order_points` to get the corners in order.
 - It calculates the document's width and height by computing Euclidean distances between points (e.g., bottom-right to bottom-left for width).

- The maximum width and height are used to define the output image size.
- A destination array (dst) defines the new image's corners (top-left at (0,0), etc.).
- `cv2.getPerspectiveTransform` computes a transformation matrix *M* to map the document's corners to the destination rectangle.
- `cv2.warpPerspective` applies the transformation to the original image, producing a flattened view.
- **Why It Matters:** This step corrects perspective distortions, making the document appear as if scanned flat.

Step 10: Enhancing the Scanned Document

- **What's Happening:** The warped image is converted to grayscale and thresholded to create a clean, black-and-white scanned effect.
- **Python Code:**

```
warped_gray = cv2.cvtColor(warped_image, cv2.COLOR_BGR2GRAY)
warped_threshold = threshold_local(warped_gray, 11, offset=10, method="gaussian")
warped_image = (warped_gray > warped_threshold).astype("uint8") * 255
cv2.imshow("Scanned Document", warped_image)
cv2.waitKey(0)
```

- **Explanation:**
 - `cv2.cvtColor` converts the warped image to grayscale.
 - `threshold_local` applies adaptive thresholding (using a 11x11 pixel neighborhood and Gaussian method) to enhance text by separating it from the background.
 - The result is a binary image (black text on white background) achieved by comparing pixel values to the threshold and scaling to 0 or 255.
 - The final image is displayed.
- **Why It Matters:** Thresholding enhances readability, mimicking the look of a scanned document.

Step 11: Cleaning Up

- **What's Happening:** The code closes all display windows.
- **Python Code:**

```
cv2.destroyAllWindows()
```

- **Explanation:** This ensures all OpenCV windows are closed, freeing up system resources.
 - **Why It Matters:** Proper cleanup prevents memory leaks and keeps the program tidy.
-